

March 30, 2024

```
[1]: import pandas as pd
```

```
[8]: df = pd.read_csv("https://raw.githubusercontent.com/sakshi2k/Social_Network_Ads/
↳master/Social_Network_Ads.csv")
print(df.dtypes)
print(df.tail())
```

```
User ID          int64
Gender           object
Age             int64
EstimatedSalary  int64
Purchased        int64
dtype: object
```

	User ID	Gender	Age	EstimatedSalary	Purchased
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

```
[5]: #input data
x = df[['Age', 'EstimatedSalary']]

# output data
y = df["Purchased"]
```

```
[6]: from sklearn.preprocessing import MinMaxScaler
sc = MinMaxScaler()
x_Scaled = sc.fit_transform(x)
```

```
[16]: from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(x_Scaled, y,
↳random_state=42, test_size=0.25)
```

```
[17]: print(Y_train)
```


```
247    1
110    0
```

```
16      1
66      0
153     0
..
71      0
106     0
270     0
348     0
102     0
Name: Purchased, Length: 300, dtype: int64
```

```
[22]: from sklearn.linear_model import LogisticRegression
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[26]: sns.countplot(x=y)
```

```
[26]: <Axes: xlabel='Purchased', ylabel='count'>
```



```
[27]: y.value_counts()
```

```
[27]: 0    257
      1    143
      Name: Purchased, dtype: int64
```

```
[28]: clasi = LogisticRegression()
      clasi.fit(X_train, Y_train)
```

```
[28]: LogisticRegression()
```

```
[29]: Y_pred = clasi.predict(X_test)
      print("Shape of Y_train: ", Y_train.shape)
      print("Shape of X_train: ", X_train.shape)
```

```
Shape of Y_train: (300,)
Shape of X_train: (300, 2)
```

```
[31]: print(Y_pred)
      print("\n\n", Y_test)
```

```
[0 1 0 1 0 0 1 0 0 0 0 1 0 0 0 0 1 1 0 1 0 0 0 1 0 0 0 0 1 0 0 0 1 0 1 0 0
 0 0 0 0 0 0 0 0 1 0 0 0 0 0 1 0 0 0 0 1 0 0 0 0 0 1 0 0 0 1 1 0 0 1 0 0 0
 0 0 1 1 0 0 0 0 1 0 1 0 0 1 0 0 1 0 0 0 0 0 1 0 1 0]
```

```
209    0
280    1
33     0
210    1
93     0
..
314    0
373    1
380    0
239    1
75     1
Name: Purchased, Length: 100, dtype: int64
```

```
[32]: fig, ax = plt.subplots(figsize=(10, 6))

      # Create a scatter plot with a color bar
      scatter = ax.scatter(x['Age'], x['EstimatedSalary'], c=y, cmap='coolwarm')

      # Create a color bar
      cbar = plt.colorbar(scatter)

      # Set the color bar label
      cbar.set_label('Purchased')
```


```

# Set the title of the scatter plot
ax.set_title('Scatter plot of Age and Estimated Salary')

# Set the x and y axis labels
ax.set_xlabel('Age')
ax.set_ylabel('Estimated Salary')

# Display the scatter plot
plt.show()

```



```
[33]: pd.DataFrame(x_Scaled).describe()
```

```

[33]:

```

	0	1
count	400.000000	400.000000
mean	0.467976	0.405500
std	0.249592	0.252570
min	0.000000	0.000000
25%	0.279762	0.207407
50%	0.452381	0.407407
75%	0.666667	0.540741
max	1.000000	1.000000

```
[35]: fig, ax = plt.subplots(figsize=(10, 6))

# Create a scatter plot with a color bar
scatter = ax.scatter(x_Scaled[:,0], x_Scaled[:,1], c=y, cmap='coolwarm')

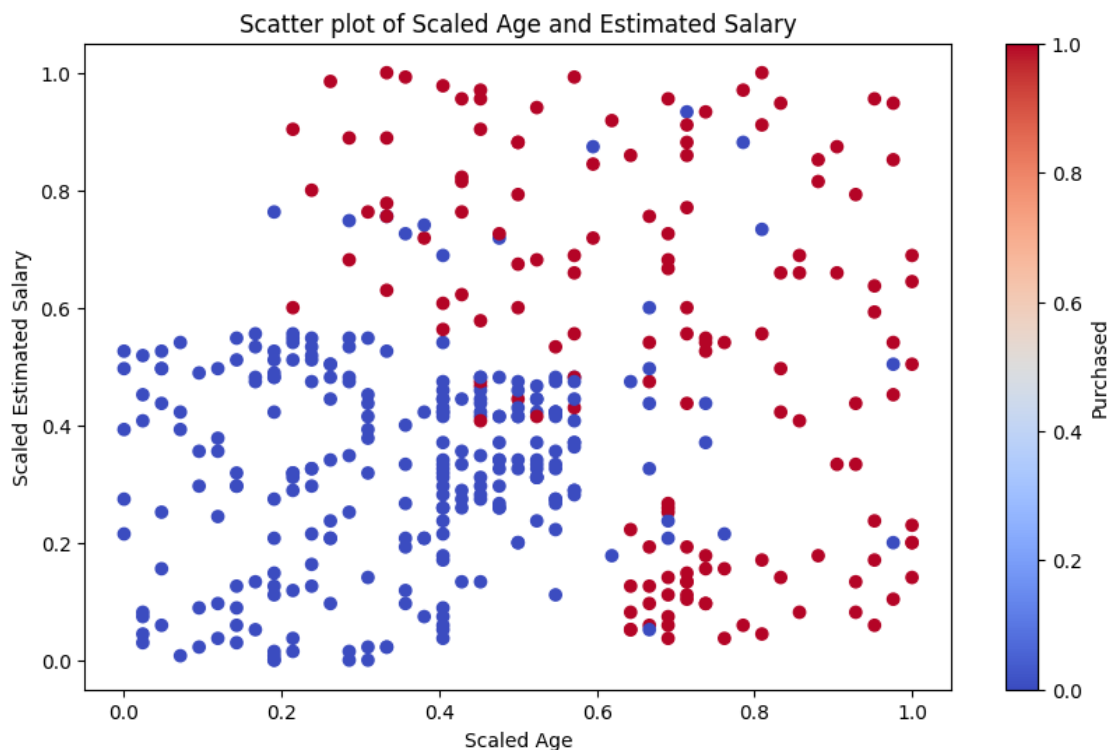
# Create a color bar
cbar = plt.colorbar(scatter)

# Set the color bar label
cbar.set_label('Purchased')

# Set the title of the scatter plot
ax.set_title('Scatter plot of Scaled Age and Estimated Salary')

# Set the x and y axis labels
ax.set_xlabel('Scaled Age')
ax.set_ylabel('Scaled Estimated Salary')

# Display the scatter plot
plt.show()
```



```
[36]: from sklearn.metrics import confusion_matrix
print(confusion_matrix(Y_test, Y_pred))
print("\n\n", Y_test.value_counts())
```

```
[[62  1]
 [12 25]]
```

```
0    63
1    37
Name: Purchased, dtype: int64
```


```
[39]: cm = confusion_matrix(Y_test, Y_pred)
fig, ax = plt.subplots(figsize=(6, 6))

# Create a heatmap of the confusion matrix with a custom color map
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', ax=ax)

# Set the title of the confusion matrix
ax.set_title('Confusion Matrix')

# Set the x and y axis labels
ax.set_xlabel('Predicted')
ax.set_ylabel('True')

# Display the confusion matrix
plt.show()
```



```
[40]: from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
print("Accuracy Score: \n", accuracy_score(Y_test,Y_pred))
print("\n\n Classification Report: \n", classification_report(Y_test,Y_pred))
```

Accuracy Score:
0.87

Classification Report:

	precision	recall	f1-score	support
0	0.84	0.98	0.91	63
1	0.96	0.68	0.79	37

accuracy			0.87	100
macro avg	0.90	0.83	0.85	100
weighted avg	0.88	0.87	0.86	100